

How to use machine learning libraries on DGX server

Zheng Lu (luz1@ornl.gov)

I will use TFGAN, a generative adversarial network library on tensorflow, as an example. Other libraries should be similar. Please replace anything in [], e.g. replace [UCAMS_id] with your 3-digit id.

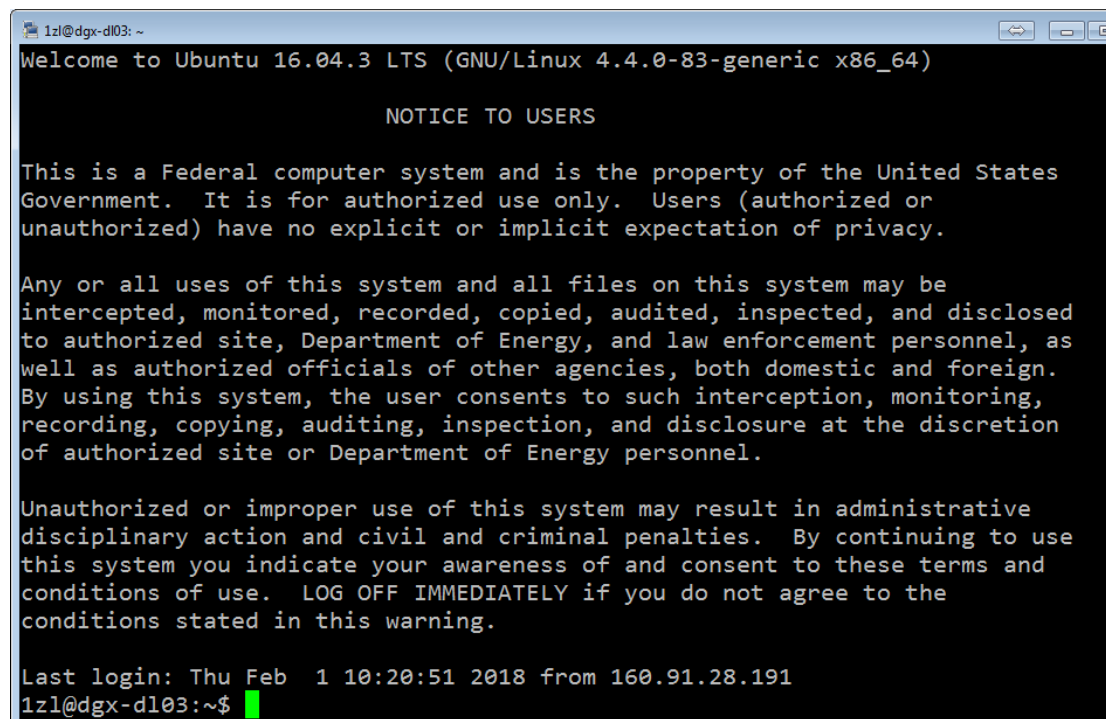
Preparation

1. Connect to DGX server using SSH.

Using your UCAMS 3-digit id and pin to login the DGX server via ssh. Windows users can use ssh softwares such as putty (<https://www.chiark.greenend.org.uk/~sgtatham/putty/>).

```
ssh [UCAMS_id]@dgx-dl03.ornl.gov
```

If you successfully login the system, you should see “Notice to users” and last login information.



```
1zl@dgx-dl03: ~  
Welcome to Ubuntu 16.04.3 LTS (GNU/Linux 4.4.0-83-generic x86_64)  
  
NOTICE TO USERS  
  
This is a Federal computer system and is the property of the United States  
Government. It is for authorized use only. Users (authorized or  
unauthorized) have no explicit or implicit expectation of privacy.  
  
Any or all uses of this system and all files on this system may be  
intercepted, monitored, recorded, copied, audited, inspected, and disclosed  
to authorized site, Department of Energy, and law enforcement personnel, as  
well as authorized officials of other agencies, both domestic and foreign.  
By using this system, the user consents to such interception, monitoring,  
recording, copying, auditing, inspection, and disclosure at the discretion  
of authorized site or Department of Energy personnel.  
  
Unauthorized or improper use of this system may result in administrative  
disciplinary action and civil and criminal penalties. By continuing to use  
this system you indicate your awareness of and consent to these terms and  
conditions of use. LOG OFF IMMEDIATELY if you do not agree to the  
conditions stated in this warning.  
  
Last login: Thu Feb  1 10:20:51 2018 from 160.91.28.191  
1zl@dgx-dl03:~$
```

2. Check your sudo privilege.

To use GPU resource on DGX server, you must have sudo privilege to run nvidia-docker. You can use the following command to check it.

```
sudo -l
```

```
1zl@dgx-dl03:~$ sudo -l
Matching Defaults entries for 1zl on dgx-dl03:
    env_reset, mail_badpass,
    secure_path=/usr/local/sbin\:/usr/local/bin\:/usr/sbin\:/usr/bin\:/sbin\:/snap/bin

User 1zl may run the following commands on dgx-dl03:
    (root) NOPASSWD: /usr/bin/docker, /usr/bin/nvidia-docker,
    /usr/bin/nvidia-smi
```

Run nvidia dockers

We use nvidia dockers on DGX server to run various machine learning tasks. We have several docker images pre-built on our DGX system. A **nvidia docker image** usually contains a machine learning library, e.g. tensorflow and several required software to run that library, e.g. python. When creating a **container** based on a docker image, you can think of all the libraries and softwares in that docker image are installed for you in a virtual environment.

1. Show Nvidia docker images.

You may want to see what docker images do we have on DGX server first with the following command.

```
sudo docker images
```

```
1zl@dgx-dl03:~$ sudo docker images
REPOSITORY          TAG                 IMAGE ID            CREATED
abm_pytorch         17.12_skllearn     6fc42ce5ee00       39 hours ago
tensorflow/tensorflow 1.5.0-devel-gpu-py3 3fc7a04f7d92       5 days ago
tensorflow/tensorflow 1.5.0-devel-gpu    cfc1345e3d5c       5 days ago
tensorflow/tensorflow latest-gpu          e6f4580b979c       5 days ago
tensorflow_jbk      latest             2ab1b6278441       6 days ago
dwp_pytorch         17.12             c2d469cdeed9       3 weeks ago
pytorch             17.12_g77         bcfa3927c5c1       3 weeks ago
nvcr.io/nvidia/pytorch 17.12             5ac6ff8f9a81       2 months ago
hello-world         latest            f2a91732366c       2 months ago
nvcr.io/nvidia/cuda  9.0-cudnn7-devel-ubuntu16.04 634be617d3ed       2 months ago
nvcr.io/nvidia/pytorch 17.11            fb6782fcfd54       2 months ago
nvcr.io/nvidia/tensorflow 17.10           9b6a599f403c       3 months ago
nvcr.io/nvidia/pytorch 17.10            41a39926dad1       3 months ago
```

2. Create container

Once you have decided which docker image to use, you can create a container for that docker.

Note that container won't save your work to harddrive automatically. So you must pass a directory name to the container when creating it. Otherwise all the data created by the container will lose once you stop or kill the container.

So, you can make a directory first with following command:

```
mkdir [data_folder]
```

Now you can create the container:

```
sudo nvidia-docker run -v [data_folder]:[path_for_mapping] -it --name=[container_name]
--net=host --rm [image_name]:[tag]
```

```
1zl@dgx-dl03:~/models/research/gan/mnist$ sudo nvidia-docker run -v /home/1z
l:/data -it --name=zlu_test1 --net=host --rm tensorflow/tensorflow:1.5.0-dev
el-gpu
root@dgx-dl03:~#
```

There are a few things to notice here:

- [path_for_mapping] is the location where your [data_folder] will mapping to in the container. Make sure to store all you computation results and intermediate results in the [path_for_mapping] directory so that after you stopping or killing your container, the data is still available.
- Remember to specify a container_name so that later you know which container is yours. By default, the system will generate a random name.
- Image_name and tag are from previous section when you run “sudo docker images”
- Once you are in the container, you can see your user name is change from your 3-digit id to root.
- Do not use “exit” or “ctrl-c” in your container command line unless you want to kill your container.

3. Check environment

You can now check if the desired machine learning library has been already installed with the following command:

```
python -c 'import tensorflow as tf; print(tf.__version__)' # for Python 2
```

```
python3 -c 'import tensorflow as tf; print(tf.__version__)' # for Python 3
```

```
root@dgx-dl03:~# python -c 'import tensorflow as tf; print(tf.__version__)'
# for Python 2
1.5.0
```

This shows tensorflow has been installed already and the version is 1.5.0

If you see “no module named ...” error, this means your library has not been installed correctly.

4. Container operations

- Detach a container:
If you still have codes running in the container and just want to exit the container and let it run in background. You can run the following command:
Ctrl-p + Ctrl-q

- Stop a container:
If you have done with your container and don't need it anymore, you can simply kill it by running the following command:

- If you are still attached to the container:
`Exit` or `Ctrl-c`

- If you have detached from the container:
`sudo docker kill [container name]`

Make sure you kill the correct container and not other people's container!
Double check your container name.

You can run the following command to see all the running containers.

`sudo docker ps`

CONTAINER ID	IMAGE	NAMES
75905b27a8e7	nvcr.io/nvidia/tensorflow:17.10	zen_brahmagupta

containername

- Re-attach to a container:
You can re-attach to the container by:
`sudo docker attach [container name]`
- Open a new bash window for a container:
You can open multiple bash window for a container:
`sudo docker exec -it [container name] /bin/bash`
When exit the extra bash window, you can simply use:
`Exit`
This will not kill the container.

Using GPU in containers

There are 4 Nvidia Tesla V100 GPU in our DGX server. Normally, when you run your code in tensorflow, it will try to allocate all GPU memory of all 4 GPUs. A good way to share the server with other people is to limit your usage to a single GPU or fraction of a single GPU.

1. Select GPU to use

To specify which GPU you want to use. You can set the environment parameter "CUDA_VISIBLE_DEVICES". You can set it either in bash or in your python code.

- In bash, this will affect all codes running after:

`export CUDA_VISIBLE_DEVICES="[index to GPUs]"`

```
root@dgx-d103:~# export CUDA_VISIBLE_DEVICES="1,3"
root@dgx-d103:~# echo $CUDA_VISIBLE_DEVICES
1,3
```

- In python code, this will only affect the code containing this command:

`import os`

`os.environ["CUDA_VISIBLE_DEVICES"] = "[index to GPUs]"`

2. Limit your GPU memory use:

For tensor flow, you can limit the GPU memory usage in your code:

```
# Assume that you want to allocate ~8GB out of 16GB:
gops = tf.GPUOptions(per_process_gpu_memory_fraction=0.5)
sess = tf.Session(config=tf.ConfigProto(gpu_options=gops))
```

Using TFGAN library

TFGAN is a light weight library for Generative Adversarial Networks. Open sourced by Google in December 2017. It provides simple function calls that cover the majority of GAN use-cases. You can just use the modules you want — loss, evaluation, features, training, etc. are all independent. When you use TFGAN, you'll be using the same infrastructure that many Google researchers use, and you'll have access to the cutting-edge improvements that Google develop with the library.

1. Requirements

The TFGAN library is built-in the tensorflow 1.5.0, earlier versions may not have it. You can test it by running the following command:

```
python -c "import tensorflow.contrib.gan"
```

2. TFGAN examples

The TFGAN examples and tutorials are in tensorflow models repository:

<https://github.com/tensorflow/models>

You can download the models repository with:

```
git clone https://github.com/tensorflow/models.git
```

The location of TFGAN is in models/research/gan

3. Run the MNIST example

You can run the MNIST example with the following commands:

```
cd models/research/gan/mnist
./launch_jobs.sh [gan_type] [models_location]
```

```
root@dgx-dl03:/data# cd models/research/gan/mnist
root@dgx-dl03:/data/models/research/gan/mnist# ./launch_jobs.sh unconditional /data/models
Dataset files already exist. Exiting without re-creating them.
Starting training unconditional GAN for 300 steps...
INFO:tensorflow:Summary name Generator/fully_connected/weights:0 is illegal; using Generator
INFO:tensorflow:Summary name Generator/fully_connected/BatchNorm/beta:0 is illegal; using Ge
INFO:tensorflow:Summary name Generator/fully_connected_1/weights:0 is illegal; using Generat
```

[gan_type] can be “unconditional”, “conditional” and “infogan”

[models_location] is the directory where you download model repository to, e.g. /data/models

4. Current problems with TFGAN examples

- MNIST
The launch script contains a bug. However, I have just submitted it on the github (<https://github.com/tensorflow/models/issues/3295>). It should be solved in future versions.
- CIFAR
The cifar example won't run correctly on our DGX server. The training will run for around 20 steps and then it will hang. I don't know the exact problem yet. However, the problem might be somewhere in CUDA drivers. Because when I run the program on CPU only, it runs well.